

IMPI Benchmark (ping-pong) to measure latency among MPI processes assigned on different cores

Nicole Orzan

December 21, 2017

The aim of today's exercise was to run intel MPI ping-pong benchmark among to processors within a node and try to estimate latency and bandwidth. I worked on node cn07-33. To execute the IMPI benchmark, first of all I had to load the following module:

```
module load impi-trial/5.0.1.035
```

There are 2 ways we can measure the latency: inside the same socket or between two different sockets. To do this we can use the "hwloc" command, which permits to bind a processes to a given CPU set.

To measure the latency inside the same socket we have to specify the numbers of 2 cores from the same socket, while to measure it between two sockets we have to specify the numbers of 2 cores from different sockets.

To know which cores are in socket 0 and which are in socket 1 we can use the command numactl -hardware. In this case I obtained the result:

```
available: 2 nodes (0-1)
node 0 cpus: 0 1 2 3 4 5 6 7 8 9
node 0 size: 163811 MB
node 0 free: 160034 MB
node 1 cpus: 10 11 12 13 14 15 16 17 18 19
node 1 size: 163840 MB
node 1 free: 159807 MB
node distances:
node  0  1
  0:  10  11
  1:  11  10
```

So for the latency inside socket 0 I used the command:

```
mpirun -np 2 hwloc-bind core:0 core:5 /u/shared/programs/x86_64/
intel/impi_5.0.1/bin64/IMB-MPI1 PingPong
```

While for the latency between two sockets I had to change:

```
hwloc-bind core:0 core:13
```

We can also specify the number of iterations to use writing "-iter num". I decided to use the default number of iterations (num=1000) and to perform 3 runs for each of the two cases.

For the first case (cores 0-5), the result I obtained for # bytes from 0 to 16 are the following:

#bytes	#repetitions	t[usec]	Mbytes/sec
0	1000	0.18	0.00
1	1000	0.19	5.05
2	1000	0.20	9.42
4	1000	0.20	18.83
8	1000	0.18	41.34
16	1000	0.18	82.69
0	1000	0.21	0.00
1	1000	0.22	4.40

2	1000	0.23	8.44
4	1000	0.20	18.62
8	1000	0.21	36.59
16	1000	0.22	69.19
0	1000	0.21	0.00
1	1000	0.22	4.27
2	1000	0.23	8.46
4	1000	0.22	17.62
8	1000	0.22	35.32
16	1000	0.23	67.51

For the second case (cores 0-13) instead, the results obtained are:

#bytes	#repetitions	t [usec]	Mbytes/sec
0	1000	0.57	0.00
1	1000	0.65	1.47
2	1000	0.67	2.84
4	1000	0.66	5.78
8	1000	0.66	11.50
16	1000	0.65	23.31
0	1000	0.59	0.00
1	1000	0.63	1.50
2	1000	0.63	3.04
4	1000	0.62	6.14
8	1000	0.62	12.39
16	1000	0.62	24.46
0	1000	0.56	0.00
1	1000	0.64	1.48
2	1000	0.63	3.02
4	1000	0.64	5.94
8	1000	0.63	12.05
16	1000	0.63	24.41

We can clearly say that the latency is higher when we bind the process to work on 2 different sockets. Calculating a mean from the outcomes obtained (as the average of all the 18 values, for the 2 different cases), we can say that in the case in which we have only one socket the value of the latency is 0.20 ± 0.01 Mbytes/sec, while in the other case its value is 0.62 ± 0.02 Mbytes/sec.