

Shared Memory Programming Paradigm

Ivan Girotto – igirotto@ictp.it

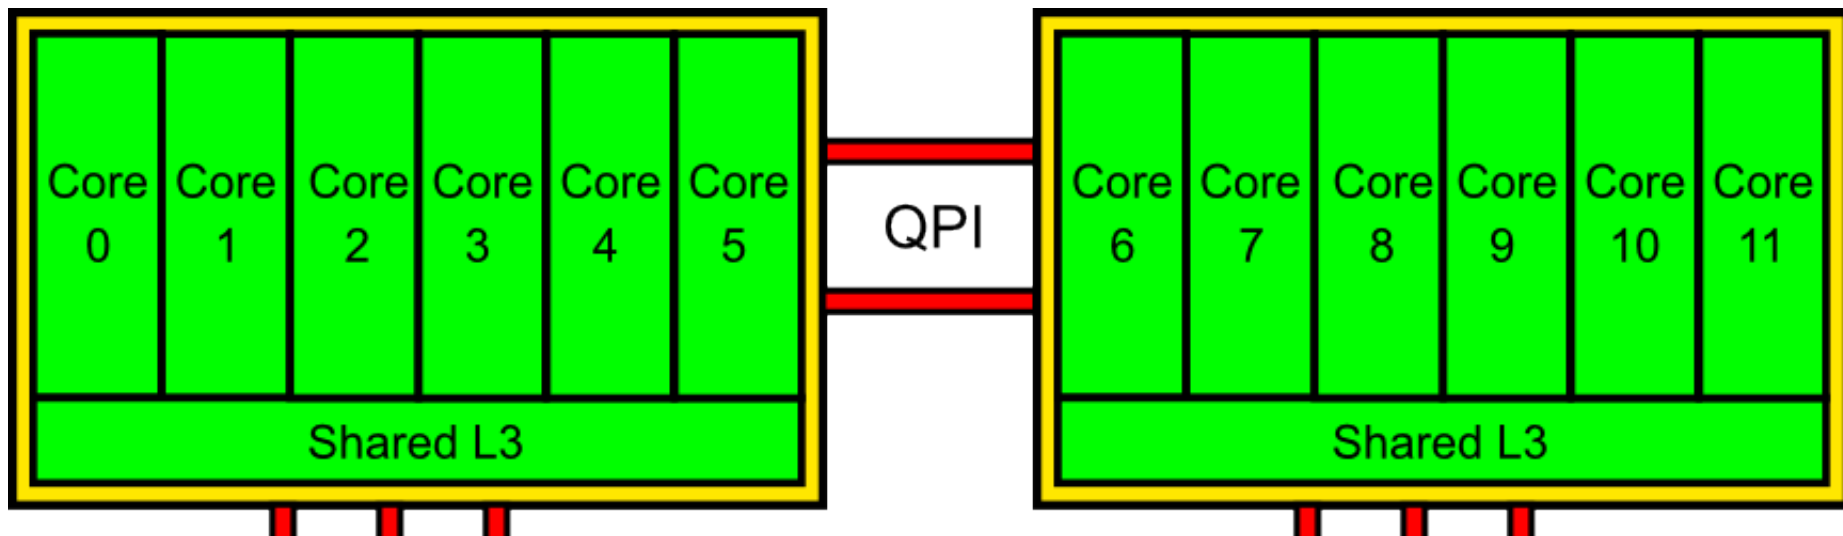
Information & Communication Technology Section (ICTS)
International Centre for Theoretical Physics (ICTP)



Scuola Internazionale Superiore
di Studi Avanzati





Multi-CPUs & Multi-cores NUMA system



Main Memory

Dual Socket (Westmere) - 24GB RAM

Processes and Threads



TEXT

compiled code (a.out)

DATA

SHARED MEMORY

STACK

System
env argv argc
auto variables for main()
auto variables for func()
<i>available for stack growth</i>
malloc.o (lib*.so)
printf.o (lib*.so)
<i>available for heap growth</i>
Heap (malloc arena)
global variables
"...%d..."
malloc.o (lib*.a)
printf.o (lib*.a)
file.o
main.o func(72,73)
crt0.o (startup routine)

High memory

← mfp – frame pointer (for main)

← stack pointer
(grows downward if func() calls another function)

↑ library functions if
dynamically linked
(usual case)

← brk point

uninitialized data (bss)

initialized data

↑ library functions if
statically linked
(not usual case)

← ra (return address)

Low memory

Stack illustrated after the call
func(72,73) called from main(),
assuming func defined by:
func(int x, int y) {
 int a;
 int b[3];
 /* no other auto variables */
}

Assumes int = long = char * of
size 4 and assumes stack at high
address and descending down.

Expanded view of the stack

Stack		Contents
Offset from current frame pointer (for func())	main() auto variables	
+12	73	y
+8	72	x
+4	ra	return address
0	mfp	caller's frame pointer
-4	garbage	a
-8	garbage	b[2]
-12	garbage	b[1]
-16	garbage	b[0]

All auto variables and parameters
are referenced via offsets from the
frame pointer.

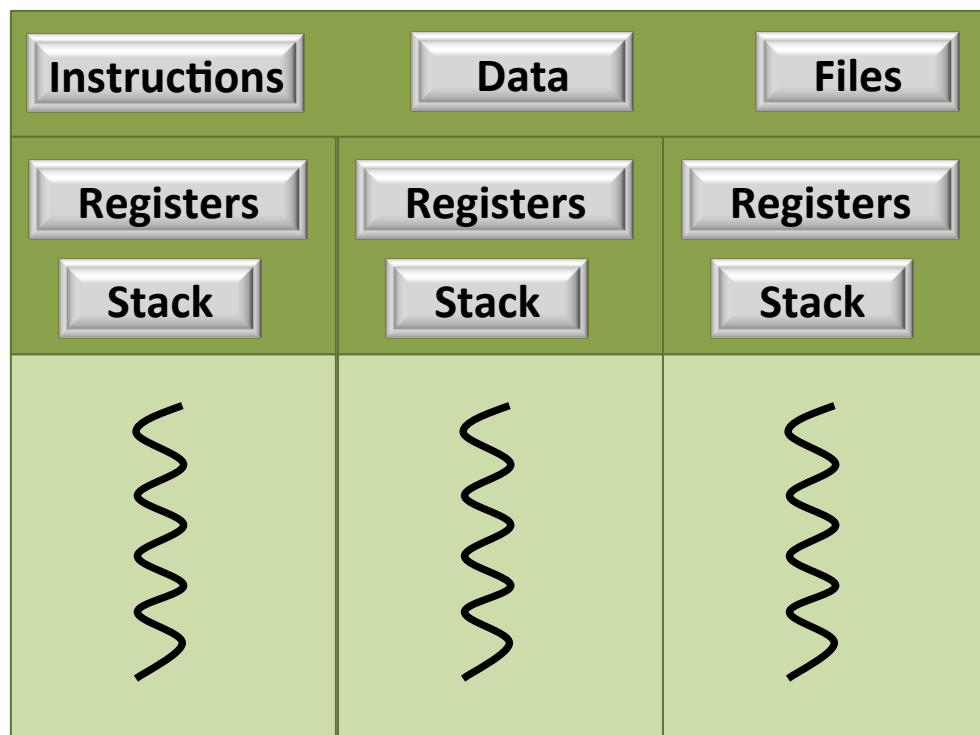
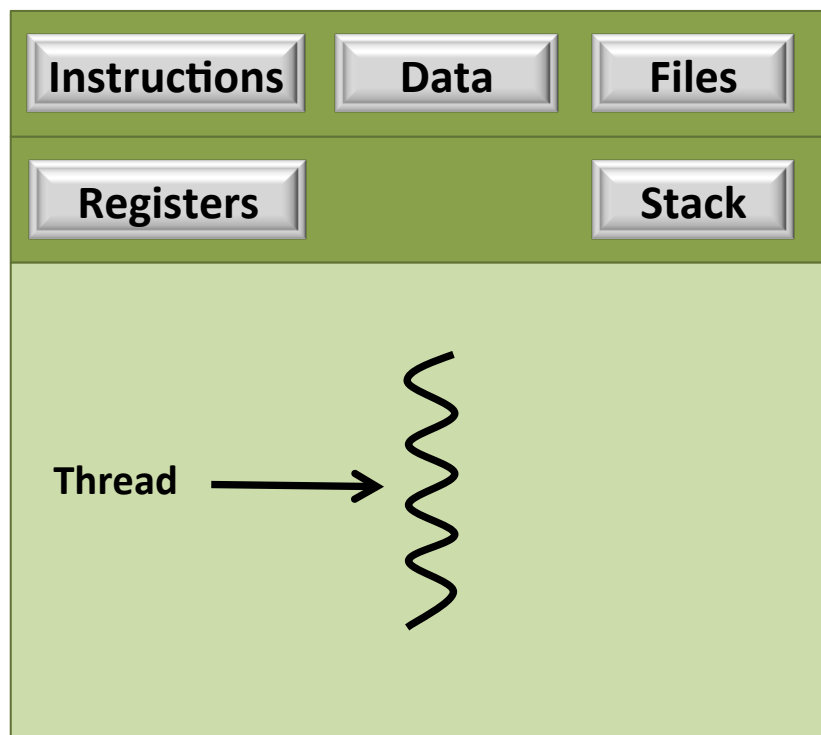
The frame pointer and stack pointer
are in registers (for fast access).

When func returns, the return value
is stored in a register. The stack pointer
is move to the y location, the code
is jumped to the return address (ra),
and the frame pointer is set to mfp
(the stored value of the caller's frame
pointer). The caller moves the return
value to the right place.

EBP

ESP

Processes and Threads



Multi-threading - Recap

- A thread is a (**lightweight**) process - an instance of a program plus its own data (private memory)
- Each thread can follow its own flow of control through a program
- Threads can share data with other threads, but also have private data
- Threads communicate with each other via the shared data.
- A *master thread* is responsible for co-ordinating the threads group

OpenMPTM

OpenMP (*Open spec. for Multi Processing*)

OpenMP **is not a computer language**

- Rather it works in conjunction with existing languages such as standard Fortran or C/C++

Application Programming Interface (API)

- that provides a **portable** model for parallel applications
- Three main components:
 - Compiler **directives**
 - **Runtime library** routines
 - **Environment variables**



OpenMP Parallelization

OpenMP is directive based

- code (can) work without them

OpenMP can be added incrementally

OpenMP only works in shared memory

- multi-socket nodes, multi-core processors


OpenMP hides the calls to a threads library

- less flexible, but much less programming

Caution: write access to shared data can easily lead to race conditions and incorrect data



OpenMP Parallelization



- Thread-based Parallelism
- Explicit Parallelism
- Fork-Join Model
- Compiler Directive Based
- Dynamic Threads

Getting Started with OpenMP

OpenMP's constructs fall into 5 categories:

- Parallel Regions
- Work sharing
- Data Environment (scope)
- Synchronization
- Runtime functions/environment variables

OpenMP is essentially the same for both Fortran and C/C++



Directives Format

A directive is a special line of source code with meaning only to certain compilers.

A directive is distinguished by a sentinel at the start of the line.

OpenMP sentinels are:

- Fortran: **!\$OMP** (or **C\$OMP** or ***\$OMP**)
- C/C++: **#pragma omp**



OpenMP: Parallel Regions

For example, to create a 4-thread parallel region:
each thread calls `foo(ID,A)` for **ID = 0** to **3**

Each thread redundantly
executes the code within
the structured block

thread-safe routine: A routine that performs
the intended function even when executed
concurrently (by more than one thread)

```
double A[1000];  
omp_set_num_threads(4);  
#pragma omp parallel  
{  
    int ID =omp_get_thread_num();  
    foo(ID,A);  
}  
printf( "All Done\n" );
```

double A[1000];

omp_set_num_threads(4);

A single copy of A is shared between all threads.

foo(0,A); foo(1,A); foo(2,A); foo(3,A);

printf("All Done\n"

Threads wait here for all threads to finish before proceeding (i.e. barrier).

How many threads?

- The number of threads in a parallel region is determined by the following factors:
 - Use of the `omp_set_num_threads()` library function
 - Setting of the `OMP_NUM_THREADS` environment variable
 - The implementation **default**
 - Threads are numbered from 0 (master thread) to N-1.

Compiling OpenMP

```
gcc -fopenmp -c my_openmp.c
```

```
gcc -fopenmp -o my_openmp.x my_openmp.o
```

```
icc -openmp -c my_openmp.c
```

```
icc -openmp -o my_openmp.x my_openmp.o
```


OpenMP runtime library

`OMP_GET_NUM_THREADS()` – returns the current # of threads.

`OMP_GET_THREAD_NUM()` - returns the id of this thread.


`OMP_SET_NUM_THREADS(n)` – set the desired # of threads.

`OMP_IN_PARALLEL()` – returns .true. if inside parallel region.

`OMP_GET_MAX_THREADS()` - returns the # of possible threads.



Memory footprint



Program



Thread 1

```
load a
add a 1
store a
```


Thread 2

```
load a
add a 1
store a
```

Private
data



Shared
data



Simple C OpenMP Program

```
#include <omp.h>
#include <stdio.h>

int main ( ) {

    printf("Starting off in the sequential world.\n");
    #pragma omp parallel
    {
        printf("Hello from thread number %d\n", omp_get_thread_num() );
    }
    printf("Back to the sequential world.\n");
    return 0;
}
```

```
PROGRAM HELLO
```

```
    INTEGER NTHREADS, TID, OMP_GET_NUM_THREADS
```

```
    INTEGER OMP_GET_THREAD_NUM
```

```
!!Fork a team of threads giving them their own copies of variables
```

```
!$OMP PARALLEL PRIVATE(NTHREADS, TID)
```

```
!!Obtain thread number
```

```
    TID = OMP_GET_THREAD_NUM()
```

```
    PRINT *, 'Hello World from thread = ', TID
```

```
!!Only master thread does this
```

```
    IF (TID .EQ. 0) THEN
```

```
        NTHREADS = OMP_GET_NUM_THREADS()
```

```
        PRINT *, 'Number of threads = ', NTHREADS
```

```
    END IF
```

```
!!All threads join master thread and disband
```

```
!$OMP END PARALLEL
```

```
END PROGRAM
```

Variable Scooping

All existing variable still exist inside a parallel region

- by default SHARED between all threads

But work sharing requires private variables

- PRIVATE clause to OMP PARALLEL directive
- Index variable of a worksharing loop
- All declared local variable within a parallel region
- The FIRSTPRIVATE clause would initialize the private instances with the contents of the shared instance

Be aware of the sharing nature of static variables



Exploiting Loop Level Parallelism

Loop level Parallelism: parallelize only loops

Easy to implement

Highly readable code

Less than optimal performance (sometimes)

Most often used

Parallel Loop Directives

Fortran do loop directive

- `!$omp do`

C\C++ for loop directive

- `#pragma omp for`

These directives do not create a team of threads but assume there has already been a team forked.

If not inside a parallel region shortcuts can be used.

- `!$omp parallel do`
- `#pragma omp parallel for`

Parallel Loop Directives /2

These are equivalent to a parallel construct followed immediately by a worksharing construct.

!\$omp parallel do

Same as

!\$omp parallel

...

!\$omp do

#pragma omp parallel for

Same as

#pragma omp parallel

...

#pragma omp for

```
integer :: N, start, len, numth, tid, i, end
double precision, dimension (N) :: a, b, c
!$OMP PARALLEL PRIVATE (start, end, len, numth, tid, i)
  numth = omp_get_num_threads()
  tid = omp_get_thread_num()
  len = N / numth
  if( tid .lt. mod( N, numth ) ) then
    len = len + 1
    start = len * tid + 1
  else
    start = len * tid + mod( N, numth ) + 1
  endif
  end = start + len - 1
  do i = start, end
    a(i) = b(i) + c(i)
  end do
!OMP END PARALLEL
```

**Not the intended
mode for OpenMP**

How is OpenMP Typically Used?

OpenMP is usually used to parallelize loops:

Split-up this loop between multiple threads

```
void main()
```

```
{
```

```
    double Res[1000];
```

```
    for(int i=0;i<1000;i++) {
```

```
        do_huge_comp(Res[i]);
```

```
    }
```

```
}
```

Sequential program

```
void main()
```

```
{
```

```
    double Res[1000];
```

```
    #pragma omp parallel for
```


```
    for(int i=0;i<1000;i++) {
```

```
        do_huge_comp(Res[i]);
```

```
    }
```

```
}
```

Parallel program



Work-Sharing Constructs

Divides the execution of the enclosed code region among the members of the team that encounter it.

Work-sharing constructs **do not** launch new threads.

No implied barrier upon entry to a work sharing construct.

However, there is an implied barrier at the end of the work sharing construct (unless **nowait** is used).



Work Sharing Constructs - example

Sequential code

```
for(i=0;I<N;i++) { a[i] = a[i] + b[i];}
```


```
#pragma omp parallel
{
    int id, i, Nthrds, istart, iend;
    id = omp_get_thread_num();
    Nthrds = omp_get_num_threads();
    istart = id * N / Nthrds;
    iend = (id+1) * N / Nthrds;
    for(i=istart;I<iend;i++) {a[i]=a[i]+b[i];}
}
```

OpenMP // Region

```
#pragma omp parallel
#pragma omp for schedule(static)
for(i=0;I<N;i++) { a[i]=a[i]+b[i];}
```

OpenMP Parallel
Region and a work-
sharing for construct

schedule(static [,chunk])




- Iterations are divided evenly among threads
- If chunk is specified, divides the work into chunk sized parcels
- If there are N threads, each thread does every Nth chunk of work.

```
!$OMP PARALLEL DO &
!$OMP SCHEDULE(STATIC,3)
```

```
DO J = 1, 36
    Work (j)
END DO
```

```
!$OMP END DO
```

schedule(dynamic [,chunk])



- Divides the workload into chunk sized parcels.
- As a thread finishes one chunk, it grabs the next available chunk.
- Default value for chunk is one.
- More overhead, but potentially better load balancing.

```
!$OMP PARALLEL DO & !
$OMPSCHEDULE(DYNAMIC,1)
```

```
DO J = 1, 36
    Work (j)
END DO
```

```
!$OMP END DO
```


The Schedule Clause SCHEDULE (type [,chunk])

The schedule clause effects how loop iterations are mapped onto threads

`schedule(static [,chunk])`

- Deal-out blocks of iterations of size “chunk” to each thread

`schedule(dynamic [,chunk])`

- Each thread grabs “chunk” iterations off a queue until all iterations have been handled

`schedule(guided [,chunk])`

- Threads dynamically grab blocks of iterations. The size of the block starts large and shrinks down to size “chunk” as the calculation proceeds

`schedule(runtime)`

- Schedule and chunk size taken from the OMP_SCHEDULE environment variable

No Wait Clauses

- No wait: if specified then threads do not synchronise at the end of the parallel loop.

For Fortran, the END DO directive is optional with NO WAIT being the default.

Note that the nowait clause is incompatible with a simple parallel region meaning that using the composite directives will not allow you to use the nowait clause.

OpenMP: Reduction(op : list)

The variables in “list” must be shared in the enclosing parallel region.

Inside a parallel or a worksharing construct:

- A local copy of each list variable is made and initialized depending on the “op” (e.g. 0 for “+”)
- pair wise “op” is updated on the local value
- Local copies are reduced into a single global copy at the end of the construct.

OpenMP: A Reduction Example

```
#include <omp.h>
#define NUM_THREADS 2
void main ()
{
    int i;
    double ZZ, func(), sum=0.0;


    omp_set_num_threads(NUM_THREADS);
    #pragma omp parallel for reduction(+:sum) private(ZZ)
    for (i=0; i< 1000; i++){
        ZZ = func(i);
        sum = sum + ZZ;
    }
}
```

Compute PI

$$\int_0^1 \frac{1}{1+x^2} dx = \arctan(x) \Big|_0^1 = \arctan(1) - \arctan(0) = \arctan(1) = \frac{\pi}{4}$$

$$\pi = 4 \int_0^1 \frac{1}{1+x^2} dx$$

Integrate, i.e determine area under function numerically using slices of $h * f(x)$ at midpoints



```
program calc_pi
  integer :: i
  real(kind=8),parameter ::PIREF = 3.141592653589793d0
  real(kind=8) :: w,x,s,pi
  integer,parameter :: n = 1000000000
  w=1.0/n
  s=0.0
  !$OMP parallel do private(i,x) reduction(+:s)
  do i=1,n
    x = w * (i - 0.5)
    s = s + 4.0 / (1.0 + x*x)
  end do
  pi = w * s
  print *, 'pi is ', pi, 'Error is ', abs(pi - PIREF)
end program calc_pi
```

if CLAUSE

We can make the parallel region directive itself conditional.

Fortran: **IF** (*scalar logical expression*)

C/C++: **if** (*scalar expression*)

```
#pragma omp parallel if (tasks > 1000)
{
    while(tasks > 0) donexttask();
}
```

SYNCHRONIZATION



Scuola Internazionale Superiore
di Studi Avanzati

OpenMP: How do Threads Interact?

OpenMP is a shared memory model.

- Threads communicate by **sharing variables**.

Unintended sharing of data can lead to **race conditions**:

- race condition: when the program's outcome changes as the threads are scheduled differently.

To control race conditions:

- Use **synchronization** to protect data conflicts.

Synchronization is expensive so:

- Change how data is stored to minimize the need for synchronization.

Note that updates to shared variables:

(e.g. $a = a + 1$)

are *not* **atomic**!

If two threads try to do this at the same time, one of the updates may get overwritten.

Program



Thread 1

```
load a
add a 1
store a
```


Thread 2

```
load a
add a 1
store a
```

Private
data



Shared
data



Barrier

Fortran

- **!\$OMP BARRIER**

C\C++

- **#pragma omp barrier**

This directive synchronises the threads in a team by causing them to wait until all of the other threads have reached this point in the code.

Implicit barriers exist after work sharing constructs. The `nowait` clause can be used to prevent this behaviour.



Add a note about single/master

Critical

Only one thread at a time can enter a **critical** section.

Example: pushing and popping a task stack

```
!$OMP PARALLEL SHARED(STACK),PRIVATE(INEXT,INEW)
```

...

```
!$OMP CRITICAL (STACKPROT)
```

```
    inext = getnext(stack)
```

```
!$OMP END CRITICAL (STACKPROT)
```

```
    call work(inext,inew)
```

```
!$OMP CRITICAL (STACKPROT)
```

```
    if (inew .gt. 0) call putnew(inew,stack)
```

```
!$OMP END CRITICAL (STACKPROT)
```

...

```
!$OMP END PARALLEL
```



Atomic

Atomic is a special case of a critical section that can be used for certain simple statements

Fortran: **!\$OMP ATOMIC**
statement

where *statement* must have one of these forms:

$x = x \text{ op } \text{expr}, \quad x = \text{expr op } x, \quad x = \text{intr} (x, \text{expr}) \text{ or}$
 $x = \text{intr}(\text{expr}, x)$

op is one of +, *, -, /, .and., .or., .eqv., Or .neqv.

intr is one of MAX, MIN, IAND, IOR or IEOB



Non Parallelizzabile

Show an example of Instruction dependency

OpenMP Tasking

- Useful to deal with unbalanced problem
- Linked lists is a good example
- Mostly applied for functional parallelism

```
#pragma omp task
```

```
!$omp task
```

Defines a task



Runtime Environment

- When a thread encounters a task construct, a new task is generated
- The moment of execution of the task is up to the runtime system
- Execution can be either immediate or delayed
- Completion of a task can be enforced through task synchronization

Task Synchronization

#pragma omp barrier

```
#pragma omp barrier
```

```
!$omp barrier
```

#pragma omp taskwait

```
#pragma omp taskwait
```

```
!$omp taskwait
```

```
my_pointer = listhead;
```

```
#pragma omp parallel  
{
```

```
    #pragma omp single  
    {
```

```
        while(my_pointer)
```

```
            #pragma omp task firstprivate(my_pointer)
```

```
            {
```

```
                (void) do_independent_work (my_pointer);
```

```
            }
```

```
            my_pointer = my_pointer->next ;
```

```
        }
```

```
    } // End of single
```

```
} // End of parallel region
```


*OpenMP Task is specified here
(executed in parallel)*



Poor Performances

Often naive approaches multi-threaded programming results in poor performances

- Modern NUMA architecture requires specific attention, specially considering multithreaded-programming
- The overhead of thread-management is not always negligible
- Reduce to minimum the critical regions
- FALSE SHARING is behind the corner
- Anything shared is a possible source of race condition (if write access)



Exercises

1. write an “hello-world” program that prints on std output how many threads are executed and the thread_ID for each thread
2. parallelize the heat equation code using OpenMP
3. parallelize the fast-transpose using OpenMP
4. parallelize the code provided in class using OpenMP

Note: perform performance analysis for the points 2-4. Write a Makefile that somehow allows to compile the serial and the OpenMP versions of the code. Instrument the code to print at runtime the number of threads, in case of a parallel version.