

# Introduction to CUDA II

... curtsey of Massimo Bernaschi (CNR - <http://www.iac.cnr.it/~massimo>) &  
John E. Stone (Univ. of Illinois at Urbana-Champaign - <http://www.ks.uiuc.edu/~johns/>)

**Ivan Girotto – [igirotto@ictp.it](mailto:igirotto@ictp.it)**

Information & Communication Technology Section (ICTS)

International Centre for Theoretical Physics (ICTP)




Scuola Internazionale Superiore  
di Studi Avanzati




# GPU Work Abstraction

- CUDA Kernels can be thought of as telling a GPU to compute all iterations of a set of nested loops concurrently
- Threads are dynamically scheduled onto hardware according to a hierarchy of thread groupings

**CPU:** Cache heavy, focused on individual thread performance




**GPU:** ALU heavy, massively parallel, throughput oriented




# Grids, Thread Blocks, Threads


**Thread blocks are scheduled onto pool of GPU SMs...**



**1-D, 2-D, 3-D thread block:**




**1-D, 2-D, or 3-D Grid of These blocks:**




# Transparent Scalability

- Hardware is free to assign blocks to any processor at any time
  - A kernel scales across any number of parallel processors




# GPU Thread Block Execution


- Thread blocks are decomposed onto hardware in **32-thread “warps”**
- Hardware execution is scheduled in units of **warps**
  - an SM can execute warps from several thread blocks
- **Warps** run in SIMD-style execution:
  - All threads execute the same instruction in lock-step
  - If one thread stalls, the entire warp stalls...
  - A branch taken by a thread has to be taken by all threads...  
(divergence is bad)



Thread blocks are multiplexed onto pool of GPU SMs...




1-D, 2-D, 3-D thread block:



# GPU Warp Branch Divergence

- Branch divergence: when not all threads take the same branch, the entire warp has to **execute both sides of the branch**
- Branch divergence issue not unique to GPUs, affects **all** SIMD hardware platforms...
- On GPUs, we get fast **hardware-based** implementation of predication/masking/etc...
- GPU blocks memory writes from disabled threads in the “if then” branch, then inverts all thread enable states and runs the “else” branch
- GPU hardware detects warp re-convergence and then runs with all threads enabled...


Thread blocks are multiplexed onto pool of GPU SMs...



# GPU Warp Branch Divergence

- Threads within the same thread block can communicate with each other in fast on-chip shared memory
- Once scheduled on an SM, thread blocks run until completion
- Because the order of thread block execution is arbitrary and blocks cannot be stopped, they cannot communicate or synchronize with other thread blocks (\*)
- (\*) Atomic memory ops are an exception wrt/communication

**Thread blocks are multiplexed onto pool of GPU SMs...**



# Execution Model

## Software


Thread




Thread  
Block

## Hardware


Scalar  
Processor



Multiprocessor



Grid



Device

Threads are executed by scalar processors

Thread blocks are executed on multiprocessors

Thread blocks **do not migrate**




Several concurrent thread blocks can reside on one multiprocessor - limited by multiprocessor resources (shared memory and register file)

A kernel is launched as a grid of thread blocks



# Memory Hierarchy

- Registers (fast up to availability)
- Local Memory: per-thread
  - Private per thread but *slow*!
  - Auto variables, register spill
- Shared Memory: per-block
  - Shared by threads of the same block
  - *Fast* inter-thread communication
- Global Memory: per-application
  - Shared by all threads
  - Inter-Grid communication



**Sequential Grids  
Execution in Time**



# Atomic Operations

- Terminology: Read-modify-write uninterruptible when *atomic*
- Many *atomic operations* on memory available with CUDA C
  - `atomicAdd()`      ▪ `atomicInc()`
  - `atomicSub()`      ▪ `atomicDec()`
  - `atomicMin()`      ▪ `atomicExch()`
  - `atomicMax()`      ▪ `atomicCAS() old == compare ? val : old`
- Predictable result when simultaneous access to memory required

# Multiblock Dot Product: `dot()`

```
__global__ void dot( int *a, int *b, int *c ) {  
    __shared__ int temp[THREADS_PER_BLOCK];  
  
    int index = threadIdx.x + blockIdx.x * blockDim.x;  
  
    temp[threadIdx.x] = a[index] * b[index];  
  
    __syncthreads();  
  
    if( 0 == threadIdx.x ) {  
        int sum = 0;  
        for( int i = 0; i < THREADS_PER_BLOCK; i++ ) sum += temp[i];  
        atomicAdd( c , sum );  
    }  
}
```

- We need to atomically add `sum` to `c` in our multiblock dot product

# Built-in Variables to manage grids and blocks

**dim3** => a new datatype defined by CUDA:

- **struct dim3 { unsigned int x, y, z };**
- three unsigned ints where any unspecified component defaults to 1.

- **dim3 gridDim;**
  - Dimensions of the grid in blocks
- **dim3 blockDim;**
  - Dimensions of the block in threads
- **dim3 blockIdx;**
  - Block index within the grid
- **dim3 threadIdx;**
  - Thread index within the block

# Bi-dimensional threads configuration: set the elements of a square matrix

```
__global__ void kernel( int *a, int dimx, int dimy ) {  
    int ix  = blockIdx.x*blockDim.x + threadIdx.x;  
    int iy  = blockIdx.y*blockDim.y + threadIdx.y;  
    int idx = iy*dimx + ix;  
  
    a[idx] = idx+1;  
}
```

```
int main() {  
    int dimx = 16;  
    int dimy = 16;  
    int num_bytes = dimx*dimy*sizeof(int);  
  
    int *d_a=0, *h_a=0; // device and host pointers  
  
    h_a = (int*)malloc(num_bytes);  
    cudaMalloc( (void**)&d_a, num_bytes );  
  
    dim3 grid, block;  
    block.x = 4;  
    block.y = 4;  
    grid.x = dimx / block.x;  
    grid.y = dimy / block.y;  
  
    kernel<<<grid, block>>>( d_a, dimx, dimy );  
  
    cudaMemcpy(h_a,d_a,num_bytes,  
               cudaMemcpyDeviceToHost);  
  
    for(int row=0; row<dimy; row++) {  
        for(int col=0; col<dimx; col++)  
            printf("%d ", h_a[row*dimx+col] );  
        printf("\n");  
    }  
  
    free( h_a );  
    cudaFree( d_a );  
    return 0;  
}
```

# GPU On-Board Global Memory

GPU arithmetic rates dwarf memory bandwidth

For Kepler K40 hardware:

~4.3 SP TFLOPS vs. ~288 GB/sec

The ratio is roughly **60 FLOPS per memory reference**  
for single-precision floating point


Peak performance achieved with **“coalesced”** memory access patterns –  
patterns that result in a single hardware memory transaction for a SIMD  
**“warp” – a contiguous group of 32 threads**

# Memory Coalescing (Oversimplified explanation)


- Threads in a warp perform a read/write operation that can be serviced in a single hardware transaction
- Rules vary slightly between hardware generations, but new GPUs are much more flexible than old ones
- If all threads in a warp read from a contiguous region that's 32 items of 4, 8, or 16 bytes in size, that's an example of a coalesced access
- Multiple threads reading the same data are handled by a hardware broadcast
- Writes are similar, but multiple writes to the same location yields undefined results

# CUDA Grid/Block/Thread Decomposition

**1-D, 2-D, or 3-D  
Computational Domain**




**1-D, 2-D, 3-D  
thread block:**



Padding arrays out to full blocks  
optimizes global memory performance  
by guaranteeing memory coalescing

**1-D, 2-D, or 3-D (SM  $\geq$  2.x)  
Grid of thread blocks:**





# CUDA Compiler: nvcc basic options

- `-arch=sm_35` → enable code for a given capability
- `-G` → enable debug for device code
- `--ptxas-options=-v` → show register and memory usage
- `--maxrregcount <N>` → limit the number of registers
- `-use_fast_math` → use fast math library
- `-O3` → Enables compiler optimization
- `-ccbin compiler_path` → use a different C compiler
- `--compiler-options` → Specify options directly to the compiler/preprocessor.

# Getting Performance From GPUs

- Don't worry (much) about counting arithmetic operations...at least until you have nothing else left to do
- GPUs provide tremendous memory bandwidth, but even so, **memory bandwidth often ends up being the performance limiter**
- Keep/reuse data in **registers** as long as possible
- The main consideration when programming GPUs is **accessing memory efficiently**, and storing operands in the **most appropriate memory system** according to data size and access pattern



# Why Does GPU Accelerate Computing?

- Highly scalable design
- Higher aggregate memory bandwidth
- Huge number of low frequency cores
- Higher aggregate computational power
- Massively parallel processors for data processing

# Why Does GPU Not Accelerate Computing?

- PCI Bus bottleneck
- Synchronization weakness
- Extremely slow serialized execution
- High complexity
  - SPMD(T) + SIMD & Memory Model
- People forget about the Amdahl's law
  - accelerating only the 50% of the original code, the expected speedup can get at most a value of 2!!

# References

- <http://www.ks.uiuc.edu/Research/gpu/>
- <http://indico.ictp.it/event/a14302/other-view?view=ictp timetable> (ICTP - SMR2760)
- <http://www.iac.rm.cnr.it/~massimo/PMC.html>
- CUDA Zone: <https://developer.nvidia.com/cuda-zone>
- CUDA by Example

